

Features Description

CRIM - per capita crime rate by town

ZN - proportion of residential land zoned for lots over 25,000 sq.ft.

INDUS - proportion of non-retail business acres per town.

CHAS - Charles River dummy variable (1 if tract bounds river; 0 otherwise)

NOX - nitric oxides concentration (parts per 10 million)

RM - average number of rooms per dwelling

AGE - proportion of owner-occupied units built prior to 1940

DIS - weighted distances to five Boston employment centres

RAD - index of accessibility to radial highways

TAX - full-value property-tax rate per 10,000

PTRATIO - pupil-teacher ratio by town

B - $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town

LSTAT - % lower status of the population

MEDV - Median value of owner-occupied homes in 1000's

Import Libraries

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
        4 import seaborn as sns
        5 import scipy.stats as stats
        6 import pylab
        7 from sklearn.model_selection import train_test_split
        8 from sklearn.preprocessing import StandardScaler
        9 from sklearn.linear_model import LinearRegression
       10 from sklearn.metrics import mean_squared_error, r2_score
       11 from sklearn.preprocessing import PolynomialFeatures
```

Load Dataset

```
In [2]: 1 raw_data = pd.read_csv('Boston_Housing.csv')
2 data = raw_data.copy()
3 print(data.shape)
4 data.head()
```


(506, 14)

```
Out[2]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33	36.2

Check Missing values

```
In [3]: 1 sns.heatmap(data.isnull(),yticklabels=False)
2 plt.show()
```




- No missing value present

Check important features

In [4]:

```
1 plt.figure(figsize=(18,8))
2 sns.heatmap(data.corr(),annot=True,cmap='viridis')
3 plt.show()
```



- As LSTAT, PTRATIO & RM are highly correlated with MEDV therefore select these 3 feature for model





In [5]:

```
1 data = data[['LSTAT', 'PTRATIO', 'RM', 'MEDV']]
```

Check presence of outliers using boxplot

In [6]:

```
1 for i in data.columns:
2     plt.figure(figsize=(14,3))
3     plt.boxplot(data[i], vert=False)
4     plt.title("Detecting outliers using Boxplot")
5     plt.xlabel(i)
6     plt.show()
```



- In MEDV value 50 repeating 16 times, this is outlier therefore removed

```
In [7]: 1 data['MEDV'].value_counts().sort_index()
```

```
Out[7]: 5.0      2
        5.6      1
        6.3      1
        7.0      2
        7.2      3
        ..
        46.7     1
        48.3     1
        48.5     1
        48.8     1
        50.0    16
Name: MEDV, Length: 229, dtype: int64
```

```
In [8]: 1 data = data[data['MEDV']!=50]
```

- In LSTAT last 2 values are outliers therefore removed

```
In [9]: 1 data['LSTAT'].value_counts().sort_index()
```

```
Out[9]: 1.98      1
        2.47      1
        2.87      1
        2.94      1
        2.98      1
        ..
        34.37     1
        34.41     1
        34.77     1
        36.98     1
        37.97     1
Name: LSTAT, Length: 442, dtype: int64
```

```
In [10]: 1 data = data[data['LSTAT']<=36]
```

- In RM last value is outlier therefore removed

```
In [11]: 1 data['RM'].value_counts().sort_index()
```

```
Out[11]: 3.561     1
         3.863     1
         4.138     1
         4.368     1
         4.628     1
         ..
         8.259     1
         8.266     1
         8.337     1
         8.398     1
         8.780     1
Name: RM, Length: 430, dtype: int64
```

```
In [12]: 1 data = data[data['RM']<=8.5]
```

```
In [13]: 1 data.shape
```


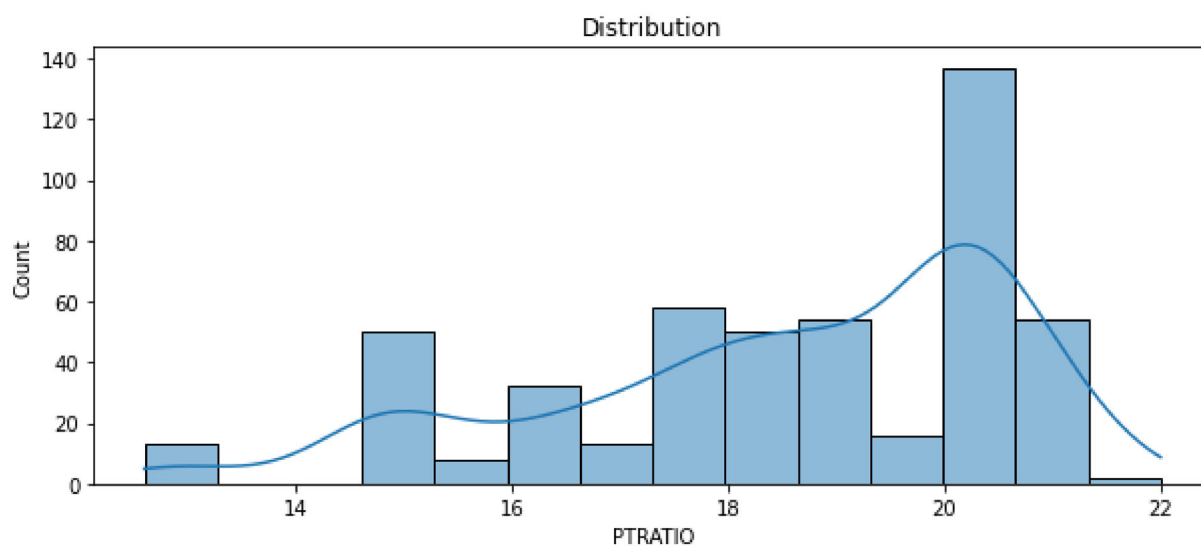

```
Out[13]: (487, 4)
```


- Total 19 data points are removed due to outliers

Now check whether features are uniformly distributed or not using histogram

In [14]:

```
1 for i in data.columns:  
2     plt.figure(figsize=(10,4))  
3     sns.histplot(data[i],kde=True)  
4     plt.title("Distribution")  
5     plt.xlabel(i)  
6     plt.show()
```





In [15]: 1 data.skew()

Out[15]: LSTAT 0.834479
PTRATIO -0.814366
RM 0.165935
MEDV 0.781963
dtype: float64

- From graph and above values, it is concluded that
- LSTAT and MEDV is right skewed
- PTRATIO is left skewed
- RM is normally distributed


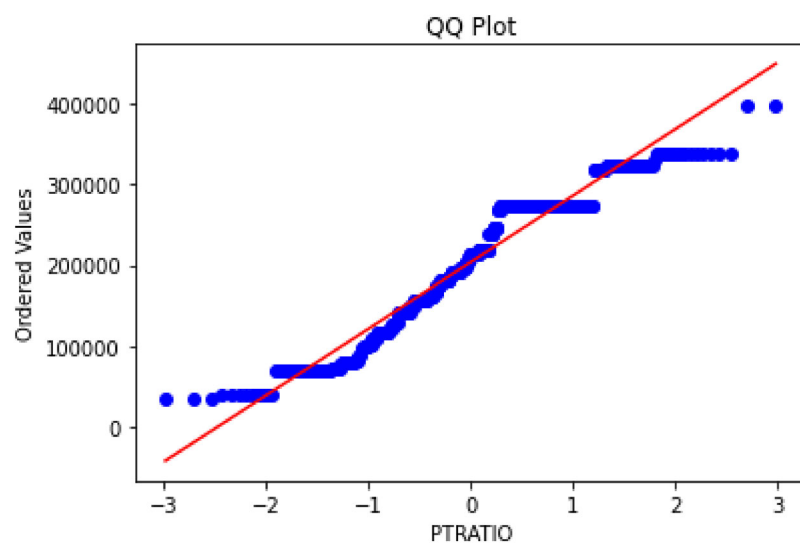

Use yeojohnson power tranformer to normalize data


In [16]: 1 data['PTRATIO'],parameters=stats.yeojohnson(data['PTRATIO'])
2 data['LSTAT'],parameters=stats.yeojohnson(data['LSTAT'])
3 data['MEDV'],parameters=stats.yeojohnson(data['MEDV'])

Check normalize data using QQ Plot

In [17]:

```
1 for i in data.columns:  
2     stats.probplot(data[i],plot=pylab)  
3     plt.title("QQ Plot")  
4     plt.xlabel(i)  
5     plt.show()
```





```
In [18]: 1 data.skew()
```

```
Out[18]: LSTAT      -0.017960
PTRATIO   -0.209212
RM         0.165935
MEDV       0.018101
dtype: float64
```

- Data is normally distributed

Seperate and splitting features and label

```
In [19]: 1 X =data[['RM', 'PTRATIO', 'LSTAT']]
2 y=data['MEDV']
```

```
In [20]: 1 X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2,random_state=
```

Use standardization technique to scale features before applying linear regression

```
In [21]: 1 X_train_stand = X_train.copy()
2 X_test_stand = X_test.copy()
3
4 # numerical features
5 num_cols = ['RM', 'PTRATIO', 'LSTAT']
6
7 # apply standardization on numerical features
8 for i in num_cols:
9
10     # fit on training data column
11     scale = StandardScaler().fit(X_train_stand[[i]])
12
13     # transform the training data column
14     X_train_stand[i] = scale.transform(X_train_stand[[i]])
15
16     # transform the testing data column
17     X_test_stand[i] = scale.transform(X_test_stand[[i]])
```

Apply linear regression model

```
In [22]: 1 lr_model = LinearRegression()
2 lr_model.fit(X_train_stand,y_train)
3 y_train_pred = lr_model.predict(X_train_stand)
4 y_test_pred = lr_model.predict(X_test_stand)
5 s1 = mean_squared_error(y_train,y_train_pred)
6 print("Mean Squared error of training set :%.2f"%s1)
7 s2 = mean_squared_error(y_test,y_test_pred)
8 print("Mean squared error of testing set: %.2f"%s2)
9
```

Mean Squared error of training set :0.304072
Mean squared error of testing set: 0.40

```
In [23]: 1 s = r2_score(y_train, y_train_pred)
2 print('R2 variance score of training set: %.2f' %s )
3 s = r2_score(y_test,y_test_pred)
4 print("R2 variance score of testing set: %.2f" %s)
5 N = y_test.size
6 p = X_train_stand.shape[1]
7 adjr2score = 1 - ((1-r2_score(y_test, y_test_pred))*(N - 1))/(N - p - 1)
8 print("Adjusted R^2 Score %.2f" % adjr2score)
```

R2 variance score of training set: 0.75
R2 variance score of testing set: 0.74
Adjusted R^2 Score 0.73

To improve r2 score ,Polynomial regression is applied

```
In [24]: 1 poly_reg = PolynomialFeatures(degree = 2)
2 X_train_poly = poly_reg.fit_transform(X_train_stand)
3 X_test_poly = poly_reg.fit_transform(X_test_stand)
4 lin_reg_2 = LinearRegression()
5 lin_reg_2.fit(X_train_poly,y_train)
```

Out[24]: LinearRegression()

```
In [25]: 1 #predicting on training data set
2 y_train_predict = lin_reg_2.predict(X_train_poly)
3 #predicting on testing data set
4 y_test_predict = lin_reg_2.predict(X_test_poly)
5 mse_train = mean_squared_error(y_train,y_train_predict)
6 r2_train = r2_score(y_train,y_train_predict)
7 print("The model performance of training set")
8 print("-----")
9 print("MSE of training set is {}".format(round(mse_train,2)))
10 print("R2 score of training set is {}".format(round(r2_train,2)))
```

The model performance of training set

MSE of training set is 0.24
R2 score of training set is 0.8

In [26]:

```
1 mse_test = mean_squared_error(y_test,y_test_predict)
2 r2_test = r2_score(y_test,y_test_predict)
3
4 print("The model performance of training set")
5 print("-----")
6 print("MSE of testing set is {}".format(round(mse_test,2)))
7 print("R2 score of testing set is {}".format(round(r2_test,2)))
```

The model performance of training set

MSE of testing set is 0.29

R2 score of testing set is 0.81

Plot of actual values vs predicted


In [27]:

```
1 plt.figure(figsize=(8,4))
2 plt.plot(y_test, y_test_predict, 'b.')
```

plt.xlabel("y_test")

plt.ylabel("y_test_predicted")

plt.show()




Difference of actual values and predicted is less

In [28]:

```
1 sns.kdeplot(y_test - y_test_predict)
```

Out[28]: <AxesSubplot:xlabel='MEDV', ylabel='Density'>



Accuracy of model is 81% & Mean square error =0.29

Create pickle file

```
In [42]: 1 import pickle
          2 file = open('model.pkl','wb')
          3 pickle.dump(lr_model,file)
```

```
In [ ]: 1
```